

Assignment No. 5

Aim :

1. Logistic Regression
2. Differentiate between Linear and Logistic Regression
3. Sigmoid Function
4. Types of LogisticRegression
5. Confusion Matrix Evaluation Metrics

Code:

```
In [43]: 1 import numpy as np
          2 import pandas as pd
          3 import matplotlib as plt
```

```
In [44]: 1 data = pd.read_csv("diabetes.csv")
          2 data
```

```
Out[44]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1
...
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

768 rows × 9 columns

```
In [45]: 1 data.head()
```

```
Out[45]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
0	6	148	72	35	0	33.6	0.627	50	1
1	1	85	66	29	0	26.6	0.351	31	0
2	8	183	64	0	0	23.3	0.672	32	1
3	1	89	66	23	94	28.1	0.167	21	0
4	0	137	40	35	168	43.1	2.288	33	1

```
In [46]: 1 data.tail()
```

```
Out[46]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
763	10	101	76	48	180	32.9	0.171	63	0
764	2	122	70	27	0	36.8	0.340	27	0
765	5	121	72	23	112	26.2	0.245	30	0
766	1	126	60	0	0	30.1	0.349	47	1
767	1	93	70	31	0	30.4	0.315	23	0

```
In [47]: 1 print("The shape of the data is: ")
        2 data.shape
```

The shape of the data is:

Out[47]: (768, 9)

```
In [48]: 1 print(data.isnull().sum())
```

```
Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction  0
Age               0
Outcome           0
dtype: int64
```

```
In [49]: 1 X = data.iloc[:,0:13]
        2 y = data.iloc[:, -1]
        3 X
        4 y
```

```
Out[49]: 0      1
        1      0
        2      1
        3      0
        4      1
        ..
       763      0
       764      0
       765      0
       766      1
       767      0
Name: Outcome, Length: 768, dtype: int64
```

```
In [47]: 1 data["Outcome"].value_counts(normalize=True)
```

```
Out[47]: Outcome
0      0.651042
1      0.348958
Name: proportion, dtype: float64
```

```
In [48]: 1 x=data.drop(["Outcome"],axis=1)
        2 y=data["Outcome"]
```

```
In [49]: 1 x
```

```
Out[49]:
```

	Pregnancies	Glucose	BloodPressure	SkinThickness	Insulin	BMI	DiabetesPedigreeFunction	Age
0	6	148	72	35	0	33.6	0.627	50
1	1	85	66	29	0	26.6	0.351	31
2	8	183	64	0	0	23.3	0.672	32
3	1	89	66	23	94	28.1	0.167	21
4	0	137	40	35	168	43.1	2.288	33
...
763	10	101	76	48	180	32.9	0.171	63
764	2	122	70	27	0	36.8	0.340	27
765	5	121	72	23	112	26.2	0.245	30
766	1	126	60	0	0	30.1	0.349	47
767	1	93	70	31	0	30.4	0.315	23

768 rows × 8 columns

```
In [50]: 1 y
```

```
Out[50]: 0      1
          1      0
          2      1
          3      0
          4      1
          ..
          763    0
          764    0
          765    0
          766    1
          767    0
          Name: Outcome, Length: 768, dtype: int64
```

```
In [50]: 1 from sklearn.model_selection import train_test_split
          2 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.20, random_state=42)
```

```
In [51]: 1 print(X_train.shape)
          2 print(X_test.shape)
          3 print(y_train.shape)
          4 print(y_test.shape)
```

```
(614, 9)
(154, 9)
(614,)
(154,)
```

```
In [51]: 1 from sklearn.preprocessing import MinMaxScaler
          2 scaler=MinMaxScaler()
          3 scaler
```

```
Out[51]: ▾ MinMaxScaler
          MinMaxScaler()
```

```
In [72]: 1 from sklearn.datasets import make_classification
          2 from sklearn.linear_model import LogisticRegression
          3 from sklearn.model_selection import train_test_split
          4 from sklearn.pipeline import make_pipeline
          5 from sklearn.preprocessing import StandardScaler
          6
          7 X, y = make_classification(random_state=42)
          8 X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42)
          9 pipe = make_pipeline(StandardScaler(), LogisticRegression())
          10 pipe.fit(X_train, y_train)
```

```
Out[72]: ▸ Pipeline
          ▸ StandardScaler
          ▸ LogisticRegression
```


```
In [78]: 1 from sklearn.linear_model import LogisticRegression
          2 logreg = LogisticRegression()
          3 logreg.fit(X_train, y_train)
          4
```

```
Out[78]: ▾ LogisticRegression
          LogisticRegression()
```

```
In [79]: 1 y_pred=logreg.predict(X_test)
```

```
In [80]: 1 from sklearn.linear_model import LinearRegression
```

```
In [81]: 1 from sklearn.preprocessing import StandardScaler
2 from sklearn.pipeline import make_pipeline
3 model = make_pipeline(StandardScaler(with_mean=False), LinearRegression())
4 model.fit(X_train, y_train)
```



```
In [82]: 1 model.score(X_test,y_test)
```

```
Out[82]: 0.6803193862233878
```

```
In [83]: 1 X_train
```

```
Out[83]: array([[ -1.06239353, -2.68317954,  0.33848384, ..., -0.35316629,
         0.32579632,  0.1943843 ],
        [ -0.79047446, -0.07873421, -1.69246463, ...,  1.09419152,
        -0.12578692,  0.05572491],
        [ -0.22096417, -0.54561186, -0.57117899, ...,  0.64084286,
        -0.28110029,  1.79768653],
        ...,
        [  0.84064355,  0.37531604, -0.96697614, ...,  0.42545756,
        0.76041466,  0.78580016],
        [  0.49403019,  0.63067073,  1.1487657 , ..., -2.84854262,
        -0.37061433,  0.77169871],
        [ -0.42018682, -0.24038388,  0.9843224 , ..., -0.99835404,
        0.23421473,  1.55050049]])
```

```
In [84]: 1 y_train
```

```
Out[84]: array([[1, 1, 1, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0,
        1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1,
        1, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1,
        1, 1, 1, 1, 1, 0, 1, 0, 1]])
```

```
In [85]: 1 X_test
```

```

-1.45518014e-01, -2.47104450e+00, -0.05565107e-01,
-2.03045386e-01,  3.71145873e-01],
[ 7.66080278e-01, -2.65461268e+00,  1.49268857e+00,
-8.26880354e-01, -2.26478892e-01,  2.25661188e+00,
-7.47211679e-01, -2.71123601e-01,  1.22693322e+00,
-1.00154076e-01,  2.36867367e+00, -9.98385439e-01,
-2.13672943e-02, -2.03673750e-01,  3.01791900e-01,
 1.17077461e+00, -8.77982587e-01, -8.03178951e-01,
 3.67365507e-01,  9.13584626e-01],
[-6.22649362e-01, -1.44723231e+00, -1.50472037e+00,
-3.70367970e-02, -4.29302225e-01,  1.25896077e+00,
-1.45755150e+00,  7.60055964e-01,  1.94607455e-01,
-7.42470586e-01,  1.27988990e+00,  4.36024464e-02,
 8.24397529e-02, -1.32002251e+00,  1.69505104e+00,
 6.63522604e-01, -6.11769091e-01, -8.31055726e-02,
-6.92420980e-01, -1.40631746e+00],
[-1.32633746e-01, -1.29826331e+00,  7.09451817e-01,
 8.47421677e-01, -5.35328186e-01,  1.25291095e+00,
-2.59546678e-01, -4.35486371e-01, -9.74529305e-01,
 1.10708067e+00,  1.09885263e+00,  6.72573701e-01,
```

```
In [86]: 1 y_test
```

```
Out[86]: array([0, 1, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1,
        1, 1, 0])
```


```
In [96]: 1 from sklearn.metrics import precision_score, ConfusionMatrixDisplay, confusion_matrix, accuracy_score
2 cm = confusion_matrix(y_test, y_pred)
3 disp = ConfusionMatrixDisplay(confusion_matrix = cm)
4 print("Confusion matrix :")
5 print(cm)
```

Confusion matrix :

```
[[15  0]
 [ 0 10]]
```

```
In [97]: 1 disp.plot()
```

Out[97]: <sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x212f7a6c650>



```
In [98]: 1 true_negative = cm[0][0]
2 false_negative = cm[1][0]
3 false_positive = cm[0][1]
4 true_positive = cm[1][1]
```

```
In [99]: 1 Accuracy = (true_positive + true_negative) / (true_positive + false_positive + true_negative + false_negative)
2 Accuracy
3 # Precision
4 Precision = true_positive / (true_positive + false_positive)
5 Precision
6 # Recall
7 Recall = true_positive / (true_positive + false_negative)
8 Recall
9 # F1 Score
10 F1_Score = 2 * (Recall * Precision) / (Recall + Precision)
11 F1_Score
```

Out[99]: 1.0

```
In [100]: 1 print("Accuracy:", Accuracy)
2 print("Confusion Matrix:")
3 print(cm)
4 print("\nClassification Report:")
5 print(classification_report(y_test, y_pred))
```

Accuracy: 1.0

Confusion Matrix:

```
[[15  0]
 [ 0 10]]
```

Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	15
1	1.00	1.00	1.00	10
accuracy			1.00	25
macro avg	1.00	1.00	1.00	25
weighted avg	1.00	1.00	1.00	25

```
In [102]: 1 Accuracy
```

Out[102]: 1.0

```
In [103]: 1 Precision
```

Out[103]: 1.0

```
In [104]: 1 Recall
```

Out[104]: 1.0

```
In [105]: 1 F1_Score
```

Out[105]: 1.0

```
In [116]: 1 from sklearn.metrics import f1_score, confusion_matrix, roc_auc_score, roc_curve
2 import matplotlib as plt
```


```
In [117]: 1 auc_score=roc_auc_score(y_test,y_pred)
```

```
In [118]: 1 fpr,tpr,thresholds=roc_curve(y_test,y_pred)
```

```
In [119]: 1 thresholds
```


Out[119]: array([2, 1, 0])

```
In [124]: 1 import matplotlib.pyplot as plt
2 plt.plot(fpr, tpr, color='orange', label='ROC')
3 plt.plot([0, 1], [0, 1], color='darkblue', linestyle='--', label='ROC curve (area = 0.73)')
4 plt.xlabel('False Positive Rate')
5 plt.ylabel('True Positive Rate')
6 plt.title('Receiver Operating Characteristic (ROC) Curve')
7 plt.legend()
8 plt.show()
```



```
In [125]: 1 import seaborn as sns
2 sns.heatmap(cm, annot=True)
```

Out[125]: <Axes: >



Name : Mohan Kadambande

Roll No. : 13212 (TECO-b1)